

Dual-IMU AHRS Simulation Report

Quaternion EKF, Error-State EKF, and Mahony Filter Comparison

Simulation study using a single and dual MARG/IMU sensor model

June 23, 2026

Contents

1	Executive Summary	3
2	Notation and Frame Convention	3
3	IMU and Magnetometer Modeling	3
3.1	Ideal specific force and magnetic field	3
3.2	Gyroscope model	4
3.3	Accelerometer model	4
3.4	Magnetometer model and calibration	4
3.5	Saturation and quantization	5
4	EKF Attitude Estimation	5
4.1	State vector	5
4.2	Prediction step	5
4.3	Measurement model	6
4.4	Correction step	6
4.5	Measurement gating	7
5	Error-State Kalman Filter Attitude Estimation	7
5.1	Nominal and error states	7
5.2	Prediction step	7
5.3	Measurement correction and error injection	8
6	Mahony AHRS Implementation	8
6.1	Predicted reference directions	8
6.2	Attitude correction error	8
6.3	Bias and gyro correction	9
7	Dual-IMU Sensor-Fusion Logic	9
7.1	Assumptions	9
7.2	Weighted measurement fusion	9
7.3	Disagreement metrics	10
8	Results	10
8.1	Main-run comparison	10
8.2	Monte Carlo comparison	11
8.3	Filter comparison	12

8.4	Gyro-bias and magnetometer behavior	13
8.5	Dual-IMU disagreement checks	14
8.6	Quaternion normalization	14
8.7	Result verdict	15
9	References	15

1 Executive Summary

The uploaded result package is internally consistent and the estimator behavior is physically reasonable. No code-level fixes are required before writing up the project. The quaternion norms remain at machine precision, the magnetometer calibration reduces direction error, and the dual-IMU disagreement metrics are nonzero and bounded, which confirms that the dual-sensor path is active rather than accidentally duplicating the single-IMU path.

The updated repository now supports three attitude estimators: a full-state quaternion EKF, an error-state Kalman filter (ESKF), and a Mahony AHRS filter. The best-performing case in the updated deterministic run is the Dual-IMU ESKF. Its roll, pitch, and yaw RMSE values are

$$0.0859^\circ, \quad 0.0687^\circ, \quad 0.1583^\circ,$$

respectively. In the repository validation Monte Carlo run with 10 runs per case, Dual-IMU ESKF gives mean roll, pitch, and yaw RMSE values of

$$0.0829^\circ, \quad 0.0658^\circ, \quad 0.1541^\circ.$$

The dual-IMU implementation improves the Kalman-filter-based estimators and the Mahony baseline. EKF and ESKF remain more accurate than Mahony for the present gain tuning. Mahony is stable and useful as a lightweight AHRS baseline, but the Kalman-filter-based implementations give better attitude accuracy and better gyro-bias behavior in this simulation.

2 Notation and Frame Convention

The navigation frame is denoted by n , and the body frame by b . The attitude quaternion is written as

$$\mathbf{q}_{nb} = \begin{bmatrix} q_w & q_x & q_y & q_z \end{bmatrix}^T,$$

where $\mathbf{q}_{nb}$ maps body-frame vectors into the navigation frame through the rotation matrix $\mathbf{R}_{nb} = \mathbf{R}(\mathbf{q}_{nb})$. Therefore,

$$\mathbf{v}^n = \mathbf{R}_{nb} \mathbf{v}^b, \quad \mathbf{v}^b = \mathbf{R}_{nb}^T \mathbf{v}^n.$$

Quaternion multiplication is denoted by $\otimes$. For an angular-rate vector ω^b , the pure quaternion form is

$$\omega_q = \begin{bmatrix} 0 & \omega_x & \omega_y & \omega_z \end{bmatrix}^T.$$

The quaternion kinematics used throughout the simulation are

$$\dot{\mathbf{q}}_{nb} = \frac{1}{2} \mathbf{q}_{nb} \otimes \omega_q.$$

After numerical integration, the quaternion is normalized to maintain

$$\|\mathbf{q}_{nb}\| = 1.$$

3 IMU and Magnetometer Modeling

3.1 Ideal specific force and magnetic field

An accelerometer does not directly measure navigation-frame acceleration. It measures specific force, which is acceleration relative to gravity. If the true navigation-frame translational

acceleration is $\mathbf{a}^n$, and gravity is $\mathbf{g}^n$, then the ideal body-frame accelerometer signal is

$$\mathbf{f}^b = \mathbf{R}_{nb}^T (\mathbf{a}^n - \mathbf{g}^n). \quad (1)$$

For an attitude-only simulation, the accelerometer is mainly used as a gravity-direction sensor. When the body is not undergoing strong external acceleration, $\mathbf{f}^b$ is dominated by the gravity direction expressed in the body frame.

The ideal magnetometer measurement is modeled as a known local magnetic field vector $\mathbf{m}^n$ rotated into the body frame:

$$\mathbf{m}^b = \mathbf{R}_{nb}^T \mathbf{m}^n. \quad (2)$$

The magnetometer gives heading information because yaw rotation changes the horizontal projection of $\mathbf{m}^b$, whereas accelerometer gravity measurements alone cannot observe yaw.

3.2 Gyroscope model

The simulated gyroscope measurement is represented as

$$\tilde{\boldsymbol{\omega}}_k^b = \mathbf{S}_g \boldsymbol{\omega}_k^b + \mathbf{b}_{g,k} + \mathbf{n}_{g,k}, \quad (3)$$

where $\mathbf{S}_g$ is the diagonal scale-factor matrix, $\mathbf{b}_{g,k}$ is the gyro bias, and $\mathbf{n}_{g,k}$ is zero-mean white noise. The gyro bias is updated as a random walk:

$$\mathbf{b}_{g,k+1} = \mathbf{b}_{g,k} + \sigma_{bg} \sqrt{\Delta t} \boldsymbol{\eta}_{g,k}, \quad \boldsymbol{\eta}_{g,k} \sim \mathcal{N}(0, \mathbf{I}). \quad (4)$$

This is a standard low-cost IMU approximation: the gyro is accurate over short time windows, but any residual bias produces attitude drift if the measurement is simply integrated.

3.3 Accelerometer model

The accelerometer measurement model is

$$\tilde{\mathbf{f}}_k^b = \mathbf{S}_a \mathbf{f}_k^b + \mathbf{b}_{a,k} + \mathbf{n}_{a,k}, \quad (5)$$

with random-walk bias

$$\mathbf{b}_{a,k+1} = \mathbf{b}_{a,k} + \sigma_{ba} \sqrt{\Delta t} \boldsymbol{\eta}_{a,k}, \quad \boldsymbol{\eta}_{a,k} \sim \mathcal{N}(0, \mathbf{I}). \quad (6)$$

For the AHRS correction step, the accelerometer is normalized:

$$\mathbf{z}_{a,k} = \frac{\tilde{\mathbf{f}}_k^b}{\|\tilde{\mathbf{f}}_k^b\|}. \quad (7)$$

Only the direction is used in the attitude correction, which makes the update less sensitive to the scale of gravity.

3.4 Magnetometer model and calibration

The raw magnetometer model is

$$\tilde{\mathbf{m}}_k^b = \mathbf{S}_m \mathbf{m}_k^b + \mathbf{b}_{m,k} + \mathbf{n}_{m,k}. \quad (8)$$

The project applies a known hard-iron and scale-factor calibration before passing the magnetometer to the AHRS:

$$\hat{\mathbf{m}}_k^b = \mathbf{S}_m^{-1} \left(\tilde{\mathbf{m}}_k^b - \mathbf{b}_m \right). \quad (9)$$

The calibrated vector is then normalized:

$$\mathbf{z}_{m,k} = \frac{\hat{\mathbf{m}}_k^b}{\|\hat{\mathbf{m}}_k^b\|}. \quad (10)$$

The results show that this calibration matters. In the single-IMU cases, the magnetic direction RMSE is reduced from 1.5570° to 0.8077° . In the dual-IMU cases, the raw fused magnetic direction RMSE is 0.7154° , and the calibrated value is 0.6221° .

3.5 Saturation and quantization

The sensor output is also bounded and quantized. For a sensor range r , the saturation operation is

$$\text{sat}(y, r) = \min(\max(y, -r), r). \quad (11)$$

For N bits, the quantization step is

$$\Delta_q = \frac{2r}{2^N}, \quad (12)$$

so the quantized output is

$$y_q = \Delta_q \text{round} \left(\frac{\text{sat}(y, r)}{\Delta_q} \right). \quad (13)$$

4 EKF Attitude Estimation

4.1 State vector

The EKF estimates attitude and gyro bias. The state vector is

$$\mathbf{x}_k = \begin{bmatrix} q_w & q_x & q_y & q_z & b_{gx} & b_{gy} & b_{gz} \end{bmatrix}^T. \quad (14)$$

The quaternion carries the attitude, and the bias state compensates slowly varying gyro drift.

4.2 Prediction step

The corrected angular rate is

$$\hat{\boldsymbol{\omega}}_k^b = \tilde{\boldsymbol{\omega}}_k^b - \hat{\mathbf{b}}_{g,k}. \quad (15)$$

A first-order quaternion integration gives

$$\mathbf{q}_{nb,k+1}^- = \text{norm} \left(\mathbf{q}_{nb,k}^+ + \frac{1}{2} \left(\mathbf{q}_{nb,k}^+ \otimes \begin{bmatrix} 0 \\ \hat{\boldsymbol{\omega}}_k^b \end{bmatrix} \right) \Delta t \right), \quad (16)$$

while the gyro bias is modeled as constant over one time step:

$$\hat{\mathbf{b}}_{g,k+1}^- = \hat{\mathbf{b}}_{g,k}^+. \quad (17)$$

The nonlinear state transition is written compactly as

$$\mathbf{x}_{k+1}^- = \mathbf{f} \left(\mathbf{x}_k^+, \tilde{\boldsymbol{\omega}}_k^b \right). \quad (18)$$

The covariance is propagated with the linearized state-transition matrix

$$\mathbf{F}_k = \left. \frac{\partial \mathbf{f}}{\partial \mathbf{x}} \right|_{\mathbf{x}_k^+, \tilde{\omega}_k^b}. \quad (19)$$

The implementation uses a numerical Jacobian for clarity. The prediction covariance update is

$$\mathbf{P}_{k+1}^- = \mathbf{F}_k \mathbf{P}_k^+ \mathbf{F}_k^T + \mathbf{Q}. \quad (20)$$

4.3 Measurement model

The EKF uses two vector observations: gravity direction from the accelerometer and magnetic-field direction from the magnetometer. The normalized measurement vector is

$$\mathbf{z}_k = \begin{bmatrix} \mathbf{z}_{a,k} \\ \mathbf{z}_{m,k} \end{bmatrix} = \begin{bmatrix} \tilde{\mathbf{f}}_k^b / \|\tilde{\mathbf{f}}_k^b\| \\ \tilde{\mathbf{m}}_k^b / \|\tilde{\mathbf{m}}_k^b\| \end{bmatrix}. \quad (21)$$

The predicted gravity and magnetic vectors in the body frame are

$$\hat{\mathbf{g}}_k^b = \mathbf{R}_{nb}(\hat{\mathbf{q}}_{nb,k})^T \mathbf{e}_g, \quad \hat{\mathbf{m}}_k^b = \mathbf{R}_{nb}(\hat{\mathbf{q}}_{nb,k})^T \mathbf{m}^n, \quad (22)$$

where $\mathbf{e}_g$ is the normalized gravity reference direction and $\mathbf{m}^n$ is the normalized navigation-frame magnetic field. The EKF measurement function is therefore

$$\mathbf{h}(\mathbf{x}_k) = \begin{bmatrix} \mathbf{R}_{nb}(\mathbf{q}_{nb,k})^T \mathbf{e}_g \\ \mathbf{R}_{nb}(\mathbf{q}_{nb,k})^T \mathbf{m}^n \end{bmatrix}. \quad (23)$$

The corresponding measurement Jacobian is

$$\mathbf{H}_k = \left. \frac{\partial \mathbf{h}}{\partial \mathbf{x}} \right|_{\mathbf{x}_k^-}. \quad (24)$$

Again, the implementation uses a numerical Jacobian. This is computationally slower than an analytic Jacobian, but it is readable and adequate for the simulation study.

4.4 Correction step

The innovation is

$$\mathbf{r}_k = \mathbf{z}_k - \mathbf{h}(\mathbf{x}_k^-). \quad (25)$$

The innovation covariance and Kalman gain are

$$\mathbf{S}_k = \mathbf{H}_k \mathbf{P}_k^- \mathbf{H}_k^T + \mathbf{R}, \quad \mathbf{K}_k = \mathbf{P}_k^- \mathbf{H}_k^T \mathbf{S}_k^{-1}. \quad (26)$$

The state update is

$$\mathbf{x}_k^+ = \mathbf{x}_k^- + \mathbf{K}_k \mathbf{r}_k. \quad (27)$$

The covariance uses the Joseph stabilized form:

$$\mathbf{P}_k^+ = (\mathbf{I} - \mathbf{K}_k \mathbf{H}_k) \mathbf{P}_k^- (\mathbf{I} - \mathbf{K}_k \mathbf{H}_k)^T + \mathbf{K}_k \mathbf{R} \mathbf{K}_k^T. \quad (28)$$

Finally, the quaternion part of the state is renormalized.

4.5 Measurement gating

The update is applied only if the measured accelerometer and magnetometer norms lie inside admissible gates:

$$\alpha_{a,\min} < \|\tilde{\mathbf{f}}_k^b\| < \alpha_{a,\max}, \quad \alpha_{m,\min} < \|\hat{\mathbf{m}}_k^b\| < \alpha_{m,\max}. \quad (29)$$

This prevents gross accelerometer or magnetometer disturbances from entering the correction step.

5 Error-State Kalman Filter Attitude Estimation

5.1 Nominal and error states

The ESKF uses the same physical sensor inputs as the EKF, but separates the attitude estimate into a nominal quaternion and a small error state. The nominal state is

$$\hat{\mathbf{x}}_k = \left\{ \hat{\mathbf{q}}_{nb,k}, \hat{\mathbf{b}}_{g,k} \right\}, \quad (30)$$

and the error state is

$$\delta \mathbf{x}_k = \left[\delta \boldsymbol{\theta}_k^T \quad \delta \mathbf{b}_{g,k}^T \right]^T. \quad (31)$$

Here $\delta \boldsymbol{\theta}$ is a minimal three-parameter attitude error. This avoids treating the four quaternion components as independent additive states and is usually cleaner for quaternion-based attitude filtering.

5.2 Prediction step

The nominal quaternion is propagated with the gyro measurement after bias compensation:

$$\hat{\boldsymbol{\omega}}_k^b = \tilde{\boldsymbol{\omega}}_k^b - \hat{\mathbf{b}}_{g,k}, \quad (32)$$

$$\hat{\mathbf{q}}_{nb,k+1}^- = \text{norm} \left(\hat{\mathbf{q}}_{nb,k}^+ + \frac{1}{2} \left(\hat{\mathbf{q}}_{nb,k}^+ \otimes \begin{bmatrix} 0 \\ \hat{\boldsymbol{\omega}}_k^b \end{bmatrix} \right) \Delta t \right). \quad (33)$$

The small-angle error dynamics are approximated by

$$\delta \boldsymbol{\theta}_{k+1} \approx \left(\mathbf{I} - [\hat{\boldsymbol{\omega}}_k^b]_{\times} \Delta t \right) \delta \boldsymbol{\theta}_k - \mathbf{I} \Delta t \delta \mathbf{b}_{g,k}, \quad (34)$$

where $[\cdot]_{\times}$ denotes the skew-symmetric cross-product matrix. Therefore, the ESKF covariance transition has the structure

$$\mathbf{F}_{\delta,k} = \begin{bmatrix} \mathbf{I} - [\hat{\boldsymbol{\omega}}_k^b]_{\times} \Delta t & -\mathbf{I} \Delta t \\ \mathbf{0} & \mathbf{I} \end{bmatrix}. \quad (35)$$

The error covariance is propagated as

$$\mathbf{P}_{\delta,k+1}^- = \mathbf{F}_{\delta,k} \mathbf{P}_{\delta,k}^+ \mathbf{F}_{\delta,k}^T + \mathbf{Q}_{\delta}. \quad (36)$$

5.3 Measurement correction and error injection

The ESKF uses the same normalized accelerometer and magnetometer vector measurements as the EKF. The predicted measurement is computed from the nominal quaternion:

$$\mathbf{h}(\hat{\mathbf{q}}_{nb,k}) = \begin{bmatrix} \mathbf{R}_{nb}(\hat{\mathbf{q}}_{nb,k})^T \mathbf{e}_g \\ \mathbf{R}_{nb}(\hat{\mathbf{q}}_{nb,k})^T \mathbf{m}^n \end{bmatrix}. \quad (37)$$

The filter estimates the small error correction

$$\delta \hat{\mathbf{x}}_k = \begin{bmatrix} \delta \hat{\boldsymbol{\theta}}_k \\ \delta \hat{\mathbf{b}}_{g,k} \end{bmatrix} = \mathbf{K}_k \left(\mathbf{z}_k - \mathbf{h}(\hat{\mathbf{q}}_{nb,k}^-) \right). \quad (38)$$

The attitude error is converted to a small quaternion:

$$\delta \mathbf{q} \approx \begin{bmatrix} 1 & \frac{1}{2} \delta \boldsymbol{\theta}^T \end{bmatrix}^T, \quad (39)$$

and injected into the nominal quaternion:

$$\hat{\mathbf{q}}_{nb,k}^+ = \text{norm} \left(\hat{\mathbf{q}}_{nb,k}^- \otimes \delta \mathbf{q} \right). \quad (40)$$

The gyro-bias estimate is corrected by

$$\hat{\mathbf{b}}_{g,k}^+ = \hat{\mathbf{b}}_{g,k}^- + \delta \hat{\mathbf{b}}_{g,k}. \quad (41)$$

After injection, the error state is reset to zero and the covariance is transformed through the reset Jacobian. This is the key practical difference between the ESKF and the full-state EKF used in Section 4.

6 Mahony AHRS Implementation

The Mahony filter is a nonlinear complementary filter on attitude. In this project it is implemented with the same external interface as the EKF and ESKF: it receives gyroscope, accelerometer, and magnetometer measurements and returns a quaternion and a gyro-bias estimate. The main difference is that Mahony uses deterministic proportional-integral feedback instead of a covariance-based Kalman update.

6.1 Predicted reference directions

At time step k , the filter predicts the body-frame gravity and magnetic directions from the current quaternion:

$$\hat{\mathbf{g}}_k^b = \mathbf{R}_{nb}(\hat{\mathbf{q}}_{nb,k})^T \mathbf{e}_g, \quad \hat{\mathbf{m}}_k^b = \mathbf{R}_{nb}(\hat{\mathbf{q}}_{nb,k})^T \mathbf{m}^n. \quad (42)$$

The measured unit vectors are

$$\mathbf{z}_{a,k} = \frac{\tilde{\mathbf{f}}_k^b}{\|\tilde{\mathbf{f}}_k^b\|}, \quad \mathbf{z}_{m,k} = \frac{\hat{\mathbf{m}}_k^b}{\|\hat{\mathbf{m}}_k^b\|}. \quad (43)$$

6.2 Attitude correction error

The attitude error is formed from vector cross-products:

$$\mathbf{e}_k = \mathbf{z}_{a,k} \times \hat{\mathbf{g}}_k^b + \mathbf{z}_{m,k} \times \hat{\mathbf{m}}_k^b. \quad (44)$$

Each cross product is small when the measured and predicted directions agree. The accelerometer term primarily corrects roll and pitch; the magnetometer term contributes yaw observability.

6.3 Bias and gyro correction

The gyro-bias estimate is updated by the integral feedback law

$$\hat{\mathbf{b}}_{g,k+1} = \hat{\mathbf{b}}_{g,k} - K_i \mathbf{e}_k \Delta t. \quad (45)$$

The corrected angular rate used for quaternion propagation is

$$\hat{\boldsymbol{\omega}}_k^b = \tilde{\boldsymbol{\omega}}_k^b - \hat{\mathbf{b}}_{g,k} + K_p \mathbf{e}_k. \quad (46)$$

The quaternion is propagated using the same kinematic equation as the EKF and ESKF:

$$\hat{\mathbf{q}}_{nb,k+1} = \text{norm} \left(\hat{\mathbf{q}}_{nb,k} + \frac{1}{2} \left(\hat{\mathbf{q}}_{nb,k} \otimes \begin{bmatrix} 0 \\ \hat{\boldsymbol{\omega}}_k^b \end{bmatrix} \right) \Delta t \right). \quad (47)$$

The gains K_p and K_i control the transient correction and bias convergence. Larger K_p gives a stronger attitude correction but can make the response noisier. Larger K_i makes bias convergence faster but can also amplify low-frequency measurement disturbances.

7 Dual-IMU Sensor-Fusion Logic

7.1 Assumptions

The current dual-IMU logic assumes both units are aligned with the same body frame and mounted without a significant lever arm. Under that assumption, both IMUs observe the same ideal angular rate, specific force, and magnetic vector. Their measurement errors differ because each unit has independent noise, bias, scale factor, and magnetometer calibration parameters.

If the sensors were physically separated by non-negligible vectors $\mathbf{r}_1$ and $\mathbf{r}_2$, the accelerometer model would need lever-arm compensation:

$$\mathbf{f}_i^b = \mathbf{f}_0^b + \dot{\boldsymbol{\omega}}^b \times \mathbf{r}_i + \boldsymbol{\omega}^b \times (\boldsymbol{\omega}^b \times \mathbf{r}_i). \quad (48)$$

That effect is not included in the present project, which is appropriate for the co-located/aligned dual-IMU simulation.

7.2 Weighted measurement fusion

The project fuses the two sensor outputs before passing them to the AHRS. For a generic three-axis measurement $\mathbf{y}$, with two independent measurements $\mathbf{y}_1$ and $\mathbf{y}_2$, each axis is fused with inverse-variance weighting:

$$\mathbf{w}_1 = \frac{1}{\boldsymbol{\sigma}_1^2}, \quad \mathbf{w}_2 = \frac{1}{\boldsymbol{\sigma}_2^2}, \quad (49)$$

where the division is elementwise. The fused measurement is

$$\mathbf{y}_f = \frac{\mathbf{w}_1 \odot \mathbf{y}_1 + \mathbf{w}_2 \odot \mathbf{y}_2}{\mathbf{w}_1 + \mathbf{w}_2}. \quad (50)$$

This is applied to gyroscope, accelerometer, and calibrated magnetometer measurements. The expected benefit is that independent white noise is reduced. If two sensors have equal variance σ^2 , then simple averaging gives

$$\sigma_f = \frac{\sigma}{\sqrt{2}}. \quad (51)$$

With unequal noise levels, the inverse-variance formula naturally gives more weight to the cleaner sensor.

7.3 Disagreement metrics

The dual-IMU health metrics are the instantaneous disagreement norms:

$$d_{a,k} = \|\tilde{\mathbf{f}}_{1,k} - \tilde{\mathbf{f}}_{2,k}\|, \quad d_{g,k} = \|\tilde{\boldsymbol{\omega}}_{1,k} - \tilde{\boldsymbol{\omega}}_{2,k}\|, \quad d_{m,k} = \|\hat{\mathbf{m}}_{1,k}^b - \hat{\mathbf{m}}_{2,k}^b\|. \quad (52)$$

The reported RMS values are

$$d_{\text{rms}} = \sqrt{\frac{1}{N} \sum_{k=1}^N d_k^2}. \quad (53)$$

The dual cases show nonzero disagreement values, confirming that the dual path is using two distinct simulated sensors.

8 Results

8.1 Main-run comparison

Table 1 summarizes the six deterministic simulation cases. The ESKF has been blended into the previous EKF and Mahony comparison, so the results now directly compare single and dual IMU modes across all three estimators. The dual ESKF gives the lowest attitude RMSE in the main run.

Table 1: Main-run attitude results for all six cases.

Case	Roll RMSE [deg]	Pitch RMSE [deg]	Yaw RMSE [deg]	Bias error norm [deg/s]
Single IMU EKF	0.2286	0.1613	0.4370	0.1112
Single IMU ESKF	0.2351	0.1657	0.4408	0.0950
Single IMU Mahony	0.2899	0.2320	0.7022	0.3828
Dual IMU EKF	0.0905	0.0695	0.1701	0.2097
Dual IMU ESKF	0.0859	0.0687	0.1583	0.1370
Dual IMU Mahony	0.1817	0.1197	0.4934	0.2874

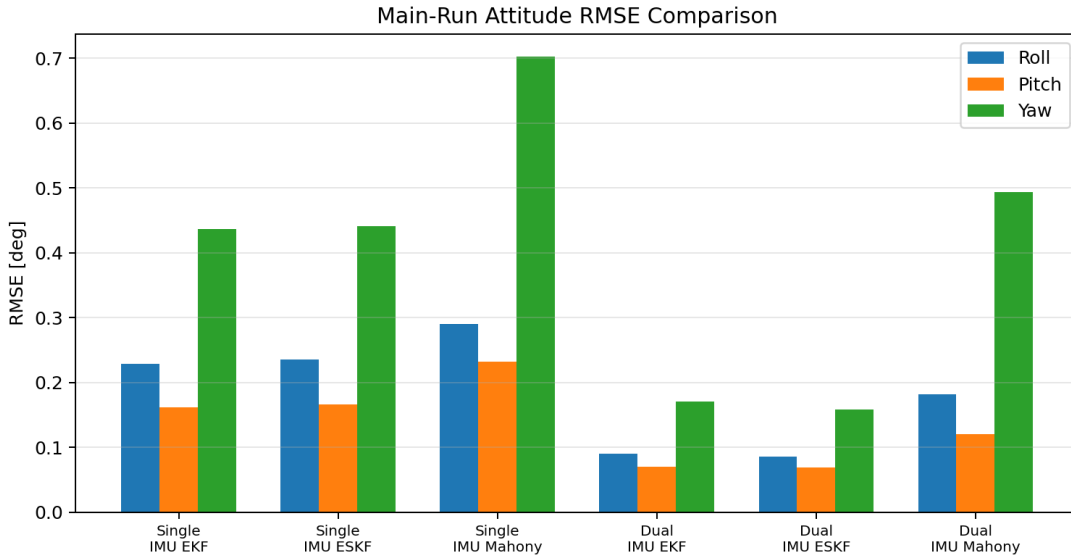


Figure 1: Main-run attitude RMSE comparison with EKF, ESKF, and Mahony included.

Relative to Single IMU EKF, Dual IMU EKF reduces main-run RMSE by approximately 60.4% in roll, 56.9% in pitch, and 61.1% in yaw. Relative to Single IMU ESKF, Dual IMU ESKF reduces main-run RMSE by approximately 63.5% in roll, 58.5% in pitch, and 64.1% in yaw. Relative to Single IMU Mahony, Dual IMU Mahony reduces main-run RMSE by approximately 37.3%, 48.4%, and 29.7%, respectively.

8.2 Monte Carlo comparison

The updated repository validation package uses 10 Monte Carlo runs per case. Table 2 reports the mean and standard deviation of attitude RMSE for all six filter/sensor combinations.

Table 2: Monte Carlo attitude RMSE over 10 runs. Values are mean $\pm$ standard deviation.

Case	Roll RMSE [deg]	Pitch RMSE [deg]	Yaw RMSE [deg]
Single IMU EKF	0.2004 ± 0.0097	0.1668 ± 0.0156	0.4292 ± 0.0601
Single IMU ESKF	0.2014 ± 0.0114	0.1675 ± 0.0171	0.4293 ± 0.0655
Single IMU Mahony	0.3358 ± 0.0675	0.3714 ± 0.1044	1.5167 ± 0.4759
Dual IMU EKF	0.0853 ± 0.0043	0.0686 ± 0.0041	0.1614 ± 0.0159
Dual IMU ESKF	0.0829 ± 0.0048	0.0658 ± 0.0050	0.1541 ± 0.0169
Dual IMU Mahony	0.2177 ± 0.0371	0.1889 ± 0.0346	0.8145 ± 0.2672

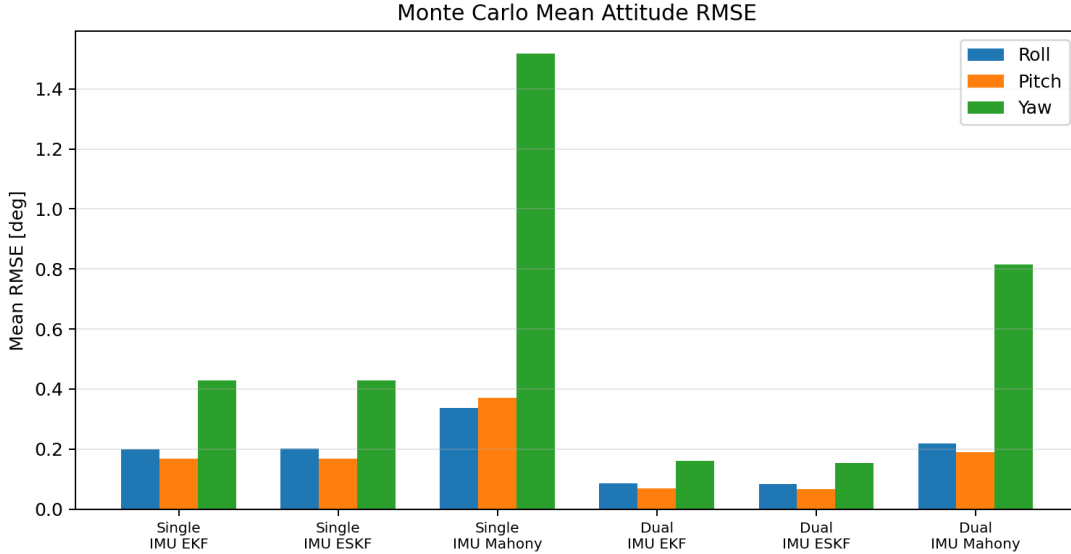


Figure 2: Monte Carlo mean attitude RMSE with EKF, ESKF, and Mahony included.

The Monte Carlo data confirms the main-run behavior. Dual-IMU fusion substantially reduces the attitude error for both Kalman-filter variants. In the ESKF case, moving from a single IMU to dual IMUs reduces mean RMSE by approximately 58.9% in roll, 60.7% in pitch, and 64.1% in yaw. Dual ESKF also gives a small but consistent improvement over Dual EKF in this validation run.

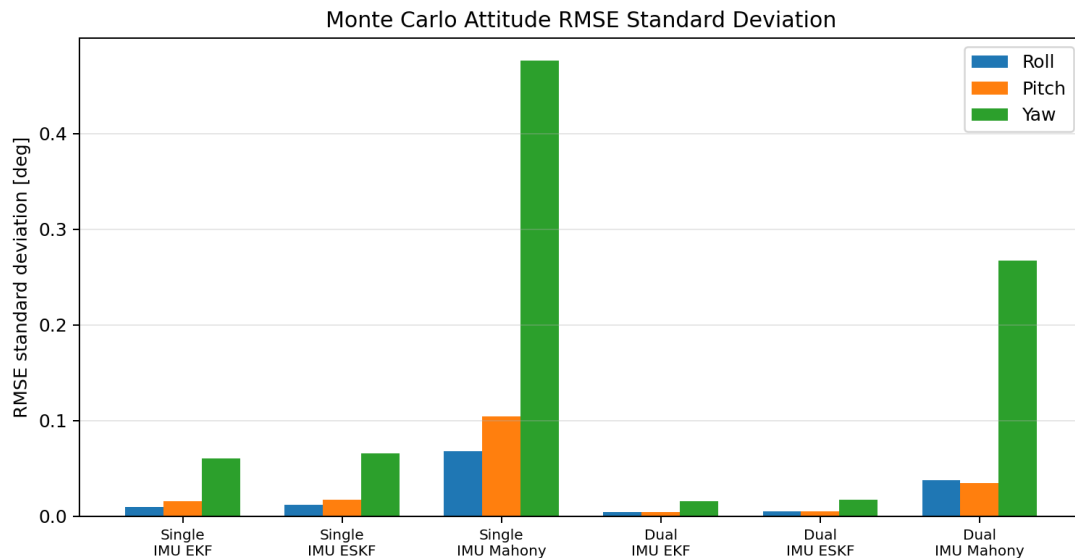


Figure 3: Monte Carlo standard deviation of attitude RMSE with ESKF included.

8.3 Filter comparison

For the dual-IMU configuration, ESKF is the best-performing filter in the updated main-run and Monte Carlo results. The full-state EKF is very close, which is expected because both filters use the same inertial and magnetic measurements. Mahony remains stable, but it has larger yaw error under the present gains.

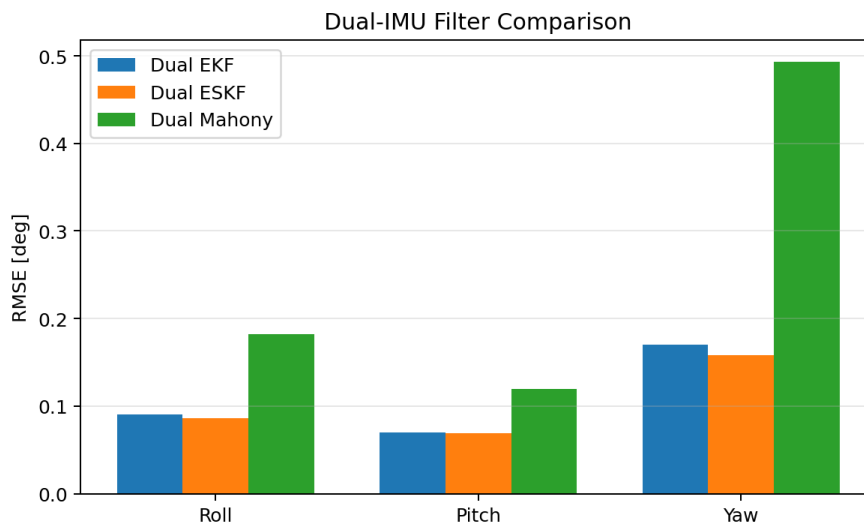


Figure 4: Dual-IMU EKF, ESKF, and Mahony attitude RMSE in the main run.

Figure 5 compares the time-domain attitude errors for the three dual-IMU filters. EKF and ESKF remain tightly bounded. Mahony is also stable, but its yaw transient is larger.

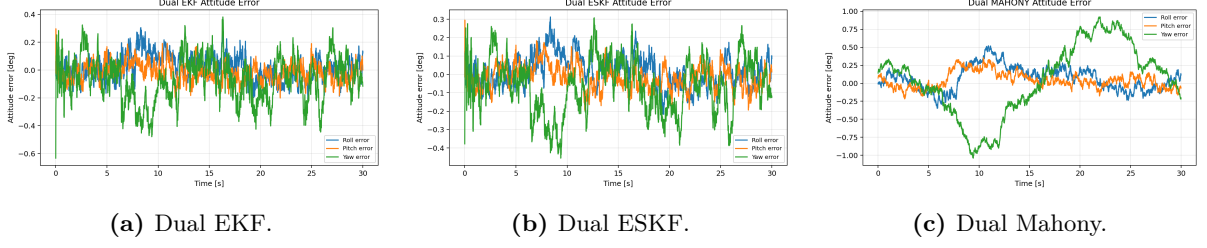


Figure 5: Time-domain attitude error for the three dual-IMU filters.

8.4 Gyro-bias and magnetometer behavior

In the deterministic main run, the final gyro-bias error norm is lowest for Single ESKF among the single-sensor cases and Dual ESKF among the dual-sensor cases. In the Monte Carlo validation, Dual ESKF also has the lowest mean final gyro-bias error norm among the dual-IMU filters.

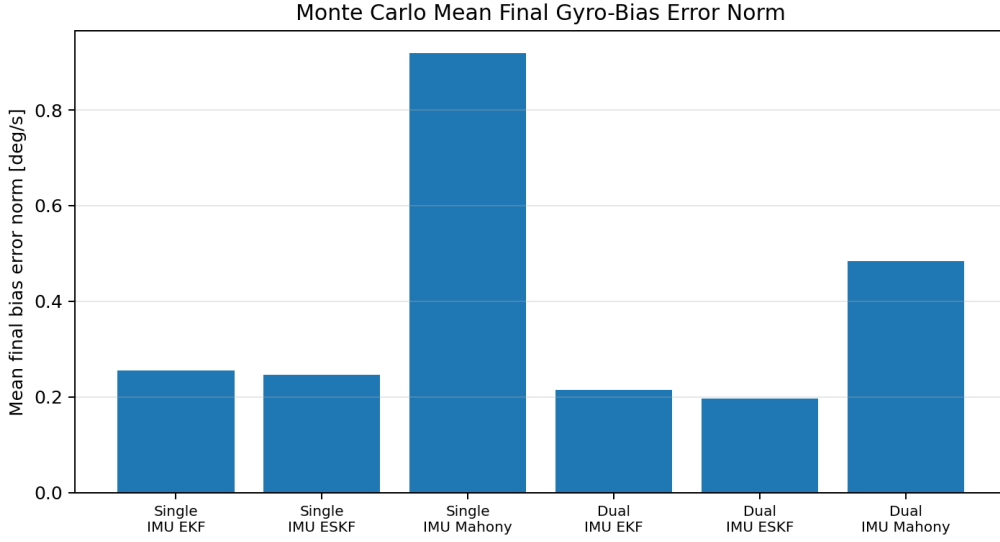


Figure 6: Monte Carlo mean final gyro-bias error norm with ESKF included.

The magnetometer direction error is improved by both calibration and dual fusion. The single-IMU calibrated magnetic direction RMSE is about 0.81° , while the dual fused calibrated value is about 0.62° . Since the sensor data is shared by filters within the same sensor mode, the magnetometer direction errors are identical across EKF, ESKF, and Mahony for the same single or dual sensor configuration.

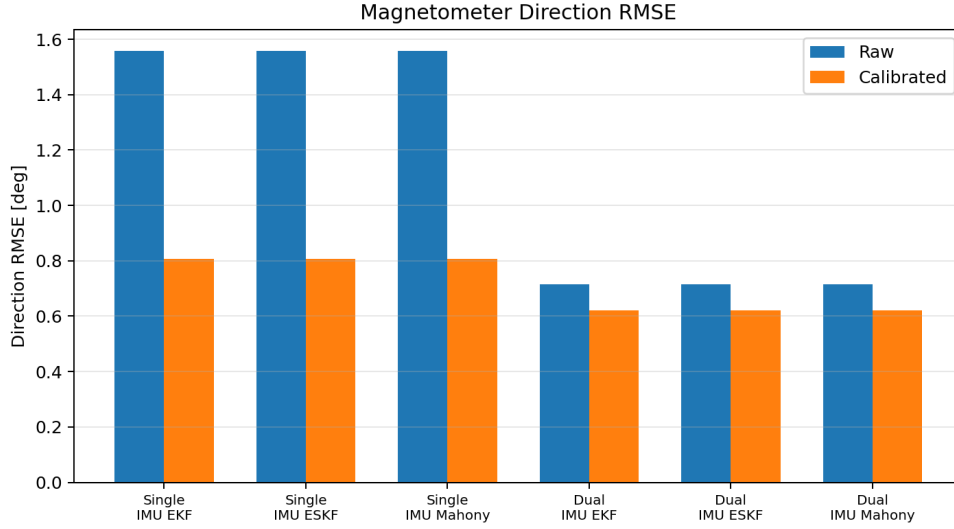


Figure 7: Magnetometer direction RMSE before and after calibration with ESKF included.

8.5 Dual-IMU disagreement checks

The dual-IMU disagreement metrics are well behaved. In the updated validation package, the mean disagreement RMS values across the dual cases are approximately

$$0.1690 \text{ m/s}^2, \quad 0.0794 \text{ rad/s}, \quad 0.0271,$$

for accelerometer, gyroscope, and calibrated magnetometer disagreement, respectively. These are nonzero, which means the two IMUs are not duplicates. They are also stable across runs, which means the measurement fusion is not erratic.

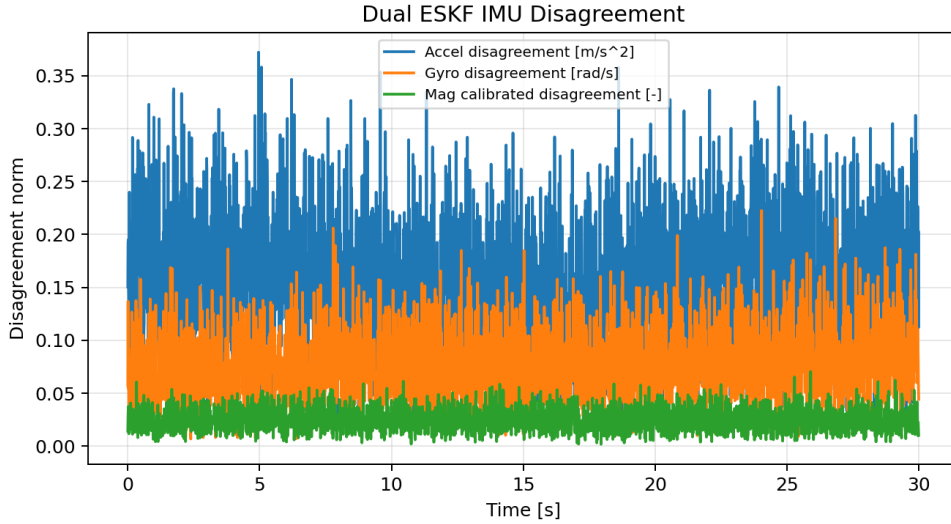


Figure 8: Dual-IMU disagreement metrics for the Dual ESKF main run.

8.6 Quaternion normalization

The true and estimated quaternion norm errors are on the order of 2.2×10^{-16} , which is essentially machine precision. This confirms that the quaternion normalization step is working for EKF, ESKF, and Mahony.

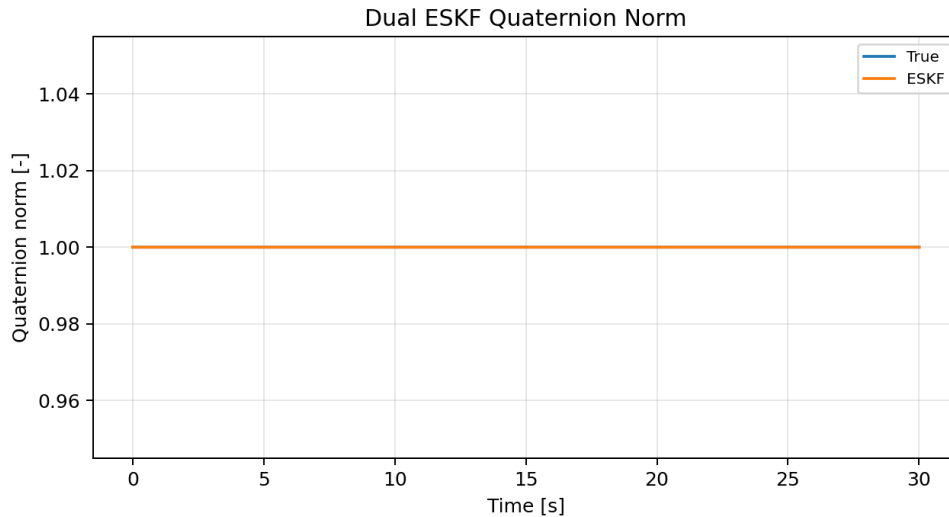


Figure 9: Quaternion norm check for the Dual-IMU ESKF case.

8.7 Result verdict

The results are good. The project demonstrates four useful points:

1. The AHRS architecture is valid: gyro propagation is stabilized by accelerometer and magnetometer vector corrections.
2. The dual-IMU measurement fusion improves attitude accuracy for EKF, ESKF, and Mahony-style filtering.
3. The ESKF is now integrated as a proper third estimator and gives the best dual-IMU result in the updated validation run.
4. The Kalman-filter-based estimators outperform Mahony for this noise model and current Mahony gain tuning.

No corrective code changes are required based on the supplied result package. If this were to be extended further, the most valuable next steps would be analytic Kalman-filter Jacobians, adaptive measurement covariance, and lever-arm compensation for spatially separated IMUs.

9 References

References

- [1] R. E. Kalman, “A new approach to linear filtering and prediction problems,” *Transactions of the ASME—Journal of Basic Engineering*, vol. 82, no. 1, pp. 35–45, 1960.
- [2] R. Mahony, T. Hamel, and J.-M. Pfimlin, “Nonlinear complementary filters on the special orthogonal group,” *IEEE Transactions on Automatic Control*, vol. 53, no. 5, pp. 1203–1218, 2008.
- [3] A. M. Sabatini, “Quaternion-based extended Kalman filter for determining orientation by inertial and magnetic sensing,” *IEEE Transactions on Biomedical Engineering*, vol. 53, no. 7, pp. 1346–1356, 2006.

- [4] P. D. Groves, *Principles of GNSS, Inertial, and Multisensor Integrated Navigation Systems*, 2nd ed. Boston, MA, USA: Artech House, 2013.
- [5] D. Titterton and J. Weston, *Strapdown Inertial Navigation Technology*, 2nd ed. Stevenage, U.K.: IET, 2004.
- [6] F. L. Markley and J. L. Crassidis, *Fundamentals of Spacecraft Attitude Determination and Control*. New York, NY, USA: Springer, 2014.
- [7] S. O. H. Madgwick, A. J. L. Harrison, and R. Vaidyanathan, "Estimation of IMU and MARG orientation using a gradient descent algorithm," in *Proc. IEEE International Conference on Rehabilitation Robotics*, 2011.
- [8] J. Solà, "Quaternion kinematics for the error-state Kalman filter," 2017.